

# Assignment 2: Test Your Neural Network on a Toy Dataset

Use this notebook to build your neural network by first implementing the following functions in the python files under assignment2/algorithms,

1. linear.py
2. relu.py
3. softmax.py
4. loss\_func.py

First you will be testing your 2 layer neural network implementation on a toy dataset.

In [65]:

```
# Setup
import matplotlib.pyplot as plt
import numpy as np

from layers.sequential import Sequential
from layers.linear import Linear
from layers.relu import ReLU
from layers.softmax import Softmax
from layers.loss_func import CrossEntropyLoss
from utils.optimizer import SGD

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots

# For auto-reloading external modules
# See http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipy
%load_ext autoreload
%autoreload 2
```

The autoreload extension is already loaded. To reload it, use:  
  %reload\_ext autoreload

We will use the class Sequential as implemented in the file assignment2/layers/sequential.py to build a layer by layer model of our neural network. Below we initialize the toy model and the toy random data that you will use to develop your implementation.

```
In [66]: # Create a small net and some toy data to check your implementations.  
# Note that we set the random seed for repeatable experiments.  
  
input_size = 4  
hidden_size = 10  
num_classes = 3 # Output  
num_inputs = 10 # N  
  
def init_toy_model():  
    np.random.seed(0)  
    l1 = Linear(input_size, hidden_size)  
    l2 = Linear(hidden_size, num_classes)  
  
    r1 = ReLU()  
    r2 = ReLU()  
    softmax = Softmax()  
    return Sequential([l1, r1, l2, softmax])  
  
def init_toy_data():  
    np.random.seed(0)  
    X = 10 * np.random.randn(num_inputs, input_size)  
    y = np.random.randint(num_classes, size=num_inputs)  
    #y = np.array([0, 1, 2, 2, 1])  
    return X, y  
  
net = init_toy_model()  
X, y = init_toy_data()
```

## Forward Pass: Compute Scores (20%)

Implement the forward functions in Linear, Relu and Softmax layers and get the output by passing our toy data X. The output must match the given output scores

```
In [67]: ┏─ scores = net.forward(X)
      ┌─ print('Your scores:')
      ┌─ print(scores)
      ┌─ print()
      ┌─ print('correct scores:')
      ┌─ correct_scores = np.asarray([[0.33333514, 0.33333826, 0.33332661],
      ┌─ [0.3333351, 0.33333828, 0.33332661],
      ┌─ [0.3333351, 0.33333828, 0.33332662],
      ┌─ [0.3333351, 0.33333828, 0.33332662],
      ┌─ [0.33333509, 0.33333829, 0.33332662],
      ┌─ [0.33333508, 0.33333829, 0.33332662],
      ┌─ [0.33333511, 0.33333828, 0.33332661],
      ┌─ [0.33333512, 0.33333827, 0.33332661],
      ┌─ [0.33333508, 0.33333829, 0.33332662],
      ┌─ [0.33333511, 0.33333828, 0.33332662]])
      ┌─ print(correct_scores)

      # The difference should be very small. We get < 1e-7
      ┌─ print('Difference between your scores and correct scores:')
      ┌─ print(np.sum(np.abs(scores - correct_scores)))
```

```
Your scores:
[[0.33333514 0.33333826 0.33332661]
 [0.3333351 0.33333828 0.33332661]
 [0.3333351 0.33333828 0.33332662]
 [0.3333351 0.33333828 0.33332662]
 [0.33333509 0.33333829 0.33332662]
 [0.33333508 0.33333829 0.33332662]
 [0.33333511 0.33333828 0.33332661]
 [0.33333512 0.33333827 0.33332661]
 [0.33333508 0.33333829 0.33332662]
 [0.33333511 0.33333828 0.33332662]]

correct scores:
[[0.33333514 0.33333826 0.33332661]
 [0.3333351 0.33333828 0.33332661]
 [0.3333351 0.33333828 0.33332662]
 [0.3333351 0.33333828 0.33332662]
 [0.33333509 0.33333829 0.33332662]
 [0.33333508 0.33333829 0.33332662]
 [0.33333511 0.33333828 0.33332661]
 [0.33333512 0.33333827 0.33332661]
 [0.33333508 0.33333829 0.33332662]
 [0.33333511 0.33333828 0.33332662]]

Difference between your scores and correct scores:
8.799388495628335e-08
```

## Forward Pass: Compute loss given the output scores from the previous step (10%)

Implement the forward function in the loss\_func.py file, and output the loss value. The loss value must match the given loss value.

```
In [68]: ┆ Loss = CrossEntropyLoss()
loss = Loss.forward(scores,y)
correct_loss = 1.098612723362578
print(loss)
# should be very small, we get < 1e-12
print('Difference between your loss and correct loss:')
print(np.sum(np.abs(loss - correct_loss)))
```

```
1.0986127233625778
Difference between your loss and correct loss:
2.220446049250313e-16
```

## Backward Pass (40%)

Implement the rest of the functions in the given files. Specifically, implement the backward function in all the 4 files as mentioned in the files. Note: No backward function in the softmax file, the gradient for softmax is jointly calculated with the cross entropy loss in the loss\_func.backward function.

You will use the chain rule to calculate gradient individually for each layer. You can assume that this calculated gradeint then is passed to the next layers in a reversed manner due to the Sequential implementation. So all you need to worry about is implementing the gradient for the current layer and multiply it will the incoming gradient (passed to the backward function as dout) to calculate the total gradient for the parameters of that layer.

```
In [69]: ┆ # No need to edit anything in this block ( 20% of the above 40% )
net.backward(Loss.backward())

gradients = []
for module in net._modules:
    for para, grad in zip(module.parameters, module.grads):
        assert grad is not None, "No Gradient"
        #Print gradients of the Linear Layer
        print(grad.shape)
        gradients.append(grad)

# Check shapes of your gradient. Note that only the Linear Layer has parameter
#(4, 10) -> Layer 1 W
#(10,) -> Layer 1 b
#(10, 3) -> Layer 2 W
#(3,) -> Layer 2 b
```

```
(4, 10)
(10,)
(10, 3)
(3,)
```

```
In [70]: # No need to edit anything in this block ( 20% of the above 40% )
# Now we check the values for these gradients. Here are the values for these
# difference, you must get difference < 1e-10
grad_w1 = np.array([[-6.24320917e-05, 3.41037180e-06, -1.69125969e-05,
                    2.41514079e-05, 3.88697976e-06, 7.63842314e-05,
                    -8.88925758e-05, 3.34909890e-05, -1.42758303e-05,
                    -4.74748560e-06],
                   [-7.16182867e-05, 4.63270039e-06, -2.20344270e-05,
                    -2.72027034e-06, 6.52903437e-07, 8.97294847e-05,
                    -1.05981609e-04, 4.15825391e-05, -2.12210745e-05,
                    3.06061658e-05],
                   [-1.69074923e-05, -8.83185056e-06, 3.10730840e-05,
                    1.23010428e-05, 5.25830316e-05, -7.82980115e-06,
                    3.02117990e-05, -3.37645284e-05, 6.17276346e-05,
                    -1.10735656e-05],
                   [-4.35902272e-05, 3.71512704e-06, -1.66837877e-05,
                    2.54069557e-06, -4.33258099e-06, 5.72310022e-05,
                    -6.94881762e-05, 2.92408329e-05, -1.89369767e-05,
                    2.01692516e-05]])
grad_b1 = np.array([-2.27150209e-06, 5.14674340e-07, -2.04284403e-06, 6.088
                     3.92085824e-06, -5.40772636e-06, 2.93354593e-06, -3.1456

grad_w2 = np.array([[ 1.28932983e-04, 1.19946731e-04, -2.48879714e-04],
                   [ 1.08784150e-04, 1.55140199e-04, -2.63924349e-04],
                   [ 6.96017544e-05, 1.42748410e-04, -2.12350164e-04],
                   [ 9.92512487e-05, 1.73257611e-04, -2.72508860e-04],
                   [ 2.05484895e-05, 4.96161144e-05, -7.01646039e-05],
                   [ 8.20539510e-05, 9.37063861e-05, -1.75760337e-04],
                   [ 2.45831715e-05, 8.74369112e-05, -1.12020083e-04],
                   [ 1.34073379e-04, 1.86253064e-04, -3.20326443e-04],
                   [ 8.86473128e-05, 2.35554414e-04, -3.24201726e-04],
                   [ 3.57433149e-05, 1.91164061e-04, -2.26907376e-04]])

grad_b2 = np.array([-0.1666649, 0.13333828, 0.03332662])

difference = np.sum(np.abs(gradients[0]-grad_w1)) + np.sum(np.abs(gradients[1]
np.sum(np.abs(gradients[3]-grad_b2))
print("Difference in Gradient values", difference)
```

Difference in Gradient values 5.665200855318841e-12

## Train the complete network on the toy data.

To train the network we will use stochastic gradient descent (SGD), we have implemented the optimizer for you. You do not implement any more functions in the python files. Below we implement the training procedure, you should get yourself familiar with the training process. Specifically looking at which functions to call and when.

Once you have implemented the method and tested various parts in the above blocks, run the code below to train a two-layer network on toy data. You should see your the training loss decrease to

In [71]: # Training Procedure

```
# Initialize the optimizer. DO NOT change any of the hyper-parameters here or
# We have implemented the SGD optimizer class for you here, which visits each
# get the gradients and optimize the respective parameters.
# You should work with the given parameters and only edit your implementation

epochs = 1000
optim = SGD(net, lr=0.1, weight_decay=0.00001)

epoch_loss = []
for epoch in range(epochs):
    # Get output scores from the network
    output_x = net(X)
    # Calculate the loss for these output scores, given the true Labe
    loss = Loss.forward(output_x, y)
    # Initialize your gradients to None in each epoch
    optim.zero_grad()
    # Make a backward pass to update the internal gradients in the La
    net.backward(Loss.backward())
    # call the step function in the optimizer to update the values of
    optim.step()
    # Append the loss at each iteration
    epoch_loss.append(loss)

    print("Epoch Loss: {:.3f}".format(epoch_loss[-1]))
```

Epoch Loss: 0.008825  
 Epoch Loss: 0.008819  
 Epoch Loss: 0.008808  
 Epoch Loss: 0.008797  
 Epoch Loss: 0.008786  
 Epoch Loss: 0.008775  
 Epoch Loss: 0.008765  
 Epoch Loss: 0.008758  
 Epoch Loss: 0.008747  
 Epoch Loss: 0.008736  
 Epoch Loss: 0.008725  
 Epoch Loss: 0.008715  
 Epoch Loss: 0.008708  
 Epoch Loss: 0.008698  
 Epoch Loss: 0.008688  
 Epoch Loss: 0.008678  
 Epoch Loss: 0.008667  
 Epoch Loss: 0.008657  
 Epoch Loss: 0.008647  
 Epoch Loss: 0.008636

In [72]: # Test your predictions. The predictions must match the labels

```
print(net.predict(X))
print(y)
```

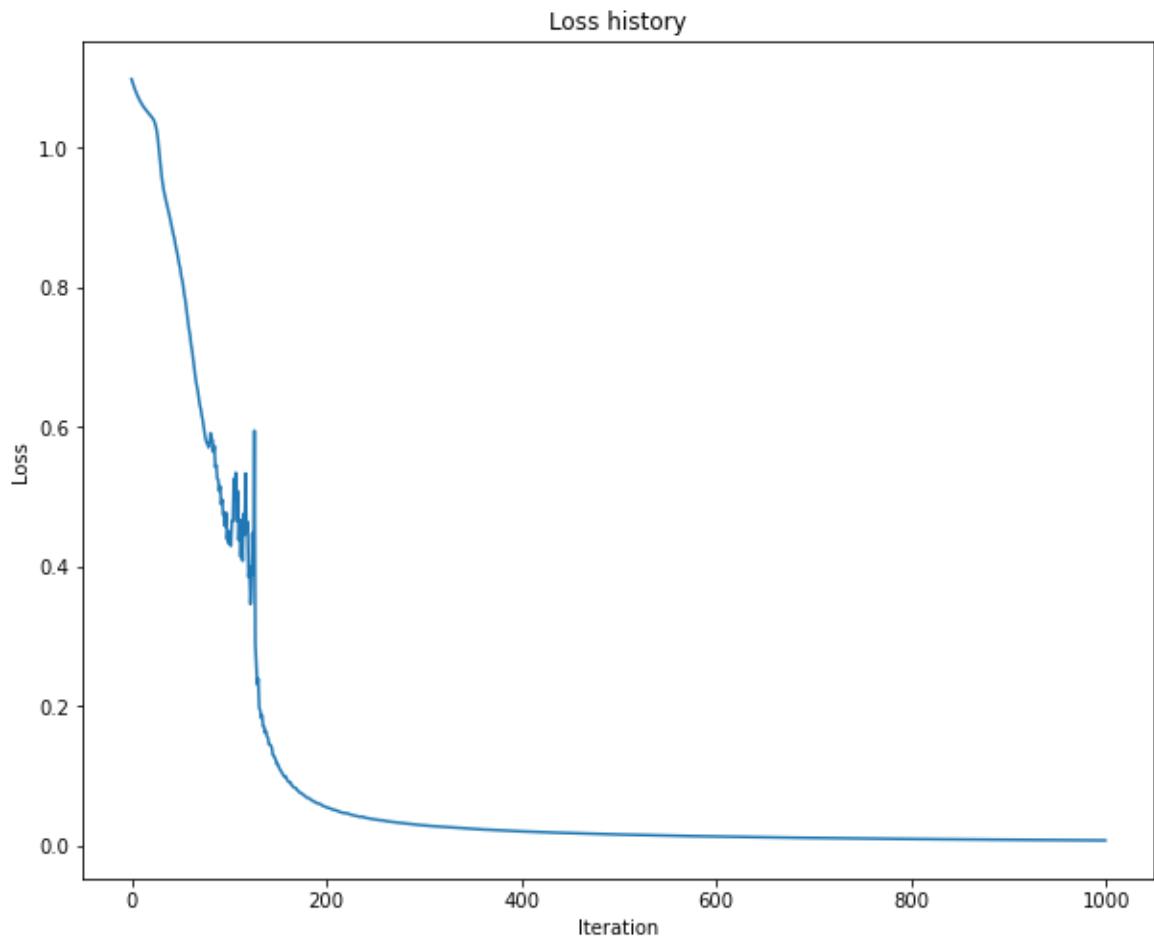
[2 1 0 1 2 0 0 2 0 0]  
[2 1 0 1 2 0 0 2 0 0]

In [73]: ► # You should be able to achieve a training loss of less than 0.02 (10%)  
print("Final training loss", epoch\_loss[-1])

Final training loss 0.007845036460819067

In [74]: ► # Plot the training Loss curve. The loss in the curve should be decreasing (2  
plt.plot(epoch\_loss)  
plt.title('Loss history')  
plt.xlabel('Iteration')  
plt.ylabel('Loss')

Out[74]: Text(0, 0.5, 'Loss')



In [ ]: ►

# Assignment 2: Train and test your model on CIFAR10 dataset

Now that you have developed and tested your model on the toy dataset set. It's time to get down and get dirty with a standard dataset such as cifar10. At this point, you will be using the provided training data to tune the hyper-parameters of your network such that it works with cifar10 for the task of multi-class classification.

Important: Recall that now we have non-linear decision boundaries, thus we do not need to do one vs all classification. We learn a single non-linear decision boundary instead. Our non-linear boundaries (thanks to relu non-linearity) will take care of differentiating between all the classes

```
In [1]: # Prepare Packages
import numpy as np
import matplotlib.pyplot as plt

from utils.data_processing import get_cifar10_data
from utils.evaluation import get_classification_accuracy

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots

# For auto-reloading external modules
# See http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

# Use a subset of CIFAR10 for the assignment
dataset = get_cifar10_data(
    subset_train = 45000,
    subset_val = 1000,
    subset_test = 10000,
)

print(dataset.keys())
print("Training Set Data Shape: ", dataset['x_train'].shape)
print("Training Set Label Shape: ", dataset['y_train'].shape)
print("Validation Set Data Shape: ", dataset['x_val'].shape)
print("Validation Set Label Shape: ", dataset['y_val'].shape)
print("Test Set Data Shape: ", dataset['x_test'].shape)
print("Test Set Label Shape: ", dataset['y_test'].shape)

dict_keys(['x_train', 'y_train', 'x_val', 'y_val', 'x_test', 'y_test'])
Training Set Data Shape: (45000, 3072)
Training Set Label Shape: (45000,)
Validation Set Data Shape: (1000, 3072)
Validation Set Label Shape: (1000,)
Test Set Data Shape: (10000, 3072)
Test Set Label Shape: (10000,)
```

```
In [2]: x_train = dataset['x_train']
y_train = dataset['y_train']
x_val = dataset['x_val']
y_val = dataset['y_val']
```

```
x_test = dataset['x_test']
y_test = dataset['y_test']
```

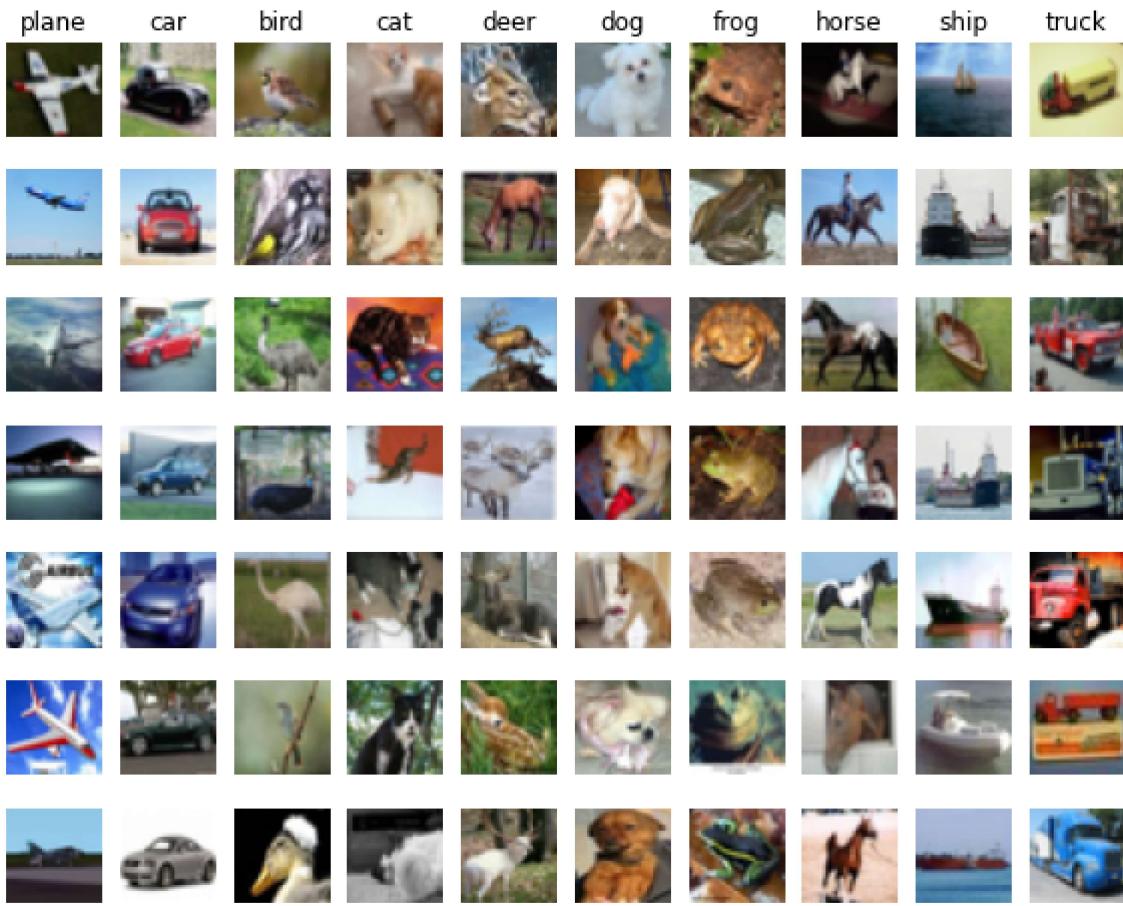
```
In [3]: # Import more utilities and the Layers you have implemented
from layers.sequential import Sequential
from layers.linear import Linear
from layers.relu import ReLU
from layers.softmax import Softmax
from layers.loss_func import CrossEntropyLoss
from utils.optimizer import SGD
from utils.dataset import DataLoader
from utils.trainer import Trainer
```

## Visualize some examples from the dataset.

```
In [4]: # We show a few examples of training images from each class.
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
samples_per_class = 7

def visualize_data(dataset, classes, samples_per_class):
    num_classes = len(classes)
    for y, cls in enumerate(classes):
        idxs = np.flatnonzero(y_train == y)
        idxs = np.random.choice(idxs, samples_per_class, replace=False)
        for i, idx in enumerate(idxs):
            plt_idx = i * num_classes + y + 1
            plt.subplot(samples_per_class, num_classes, plt_idx)
            plt.imshow(dataset[idx])
            plt.axis('off')
            if i == 0:
                plt.title(cls)
    plt.show()

visualize_data(x_train.reshape(45000, 3, 32, 32).transpose(0, 2, 3, 1), classes, sample
```



## Initialize the model

In [5]:

```
input_size = 3072
hidden_size = 50 # Hidden Layer size (Hyper-parameter)
num_classes = 10 # Output

# For a default setting we use the same model we used for the toy dataset.
# This tells you the power of a 2 layered Neural Network. Recall the Universal Approximation Theorem.
# A 2 layer neural network with non-linearities can approximate any function, given large enough hidden layers.
def init_model():
    #np.random.seed(0) # No need to fix the seed here
    l1 = Linear(input_size, hidden_size)
    l2 = Linear(hidden_size, num_classes)

    r1 = ReLU()
    softmax = Softmax()
    return Sequential([l1, r1, l2, softmax])
```

In [6]:

```
# Initialize the dataset with the dataloader class
dataset = DataLoader(x_train, y_train, x_val, y_val, x_test, y_test)
net = init_model()
optim = SGD(net, lr=0.002, weight_decay=0.00025)
loss_func = CrossEntropyLoss()
epoch = 300 # (Hyper-parameter)
batch_size = 100 # (Reduce the batch size if your computer is unable to handle it)
```

In [7]:

```
#Initialize the trainer class by passing the above modules
trainer = Trainer(dataset, optim, net, loss_func, epoch, batch_size, validate_interval=1)
```

```
In [8]: # Call the trainer function we have already implemented for you. This trains the model  
# hyper-parameters. It follows the same procedure as in the last ipython notebook you u  
train_error, validation_accuracy = trainer.train()
```

```
Epoch Average Loss: 2.302516  
Validate Acc: 0.090  
Epoch Average Loss: 2.302202  
Epoch Average Loss: 2.301113  
Epoch Average Loss: 2.297915  
Validate Acc: 0.090  
Epoch Average Loss: 2.290991  
Epoch Average Loss: 2.277926  
Epoch Average Loss: 2.259583  
Validate Acc: 0.127  
Epoch Average Loss: 2.240886  
Epoch Average Loss: 2.225274  
Epoch Average Loss: 2.213205  
Validate Acc: 0.149  
Epoch Average Loss: 2.203971  
Epoch Average Loss: 2.196885  
Epoch Average Loss: 2.191192  
Validate Acc: 0.158  
Epoch Average Loss: 2.186730  
Epoch Average Loss: 2.183068  
Epoch Average Loss: 2.180109  
Validate Acc: 0.165  
Epoch Average Loss: 2.177490  
Epoch Average Loss: 2.175474  
Epoch Average Loss: 2.173616  
Validate Acc: 0.166  
Epoch Average Loss: 2.171998  
Epoch Average Loss: 2.170910  
Epoch Average Loss: 2.169779  
Validate Acc: 0.169  
Epoch Average Loss: 2.168803  
Epoch Average Loss: 2.168057  
Epoch Average Loss: 2.167329  
Validate Acc: 0.167  
Epoch Average Loss: 2.166557  
Epoch Average Loss: 2.166109  
Epoch Average Loss: 2.165598  
Validate Acc: 0.163  
Epoch Average Loss: 2.165213  
Epoch Average Loss: 2.164571  
Epoch Average Loss: 2.164091  
Validate Acc: 0.163  
Epoch Average Loss: 2.163411  
Epoch Average Loss: 2.162704  
Epoch Average Loss: 2.161879  
Validate Acc: 0.165  
Epoch Average Loss: 2.160889  
Epoch Average Loss: 2.159469  
Epoch Average Loss: 2.157904  
Validate Acc: 0.177  
Epoch Average Loss: 2.156050  
Epoch Average Loss: 2.153642  
Epoch Average Loss: 2.150960  
Validate Acc: 0.194  
Epoch Average Loss: 2.147927  
Epoch Average Loss: 2.144576  
Epoch Average Loss: 2.141130  
Validate Acc: 0.220  
Epoch Average Loss: 2.137517  
Epoch Average Loss: 2.134063  
Epoch Average Loss: 2.130644
```

Validate Acc: 0.226  
Epoch Average Loss: 2.127413  
Epoch Average Loss: 2.124509  
Epoch Average Loss: 2.121739  
Validate Acc: 0.232  
Epoch Average Loss: 2.119247  
Epoch Average Loss: 2.116981  
Epoch Average Loss: 2.114949  
Validate Acc: 0.225  
Epoch Average Loss: 2.112995  
Epoch Average Loss: 2.110915  
Epoch Average Loss: 2.108763  
Validate Acc: 0.230  
Epoch Average Loss: 2.105745  
Epoch Average Loss: 2.099892  
Epoch Average Loss: 2.088747  
Validate Acc: 0.208  
Epoch Average Loss: 2.073367  
Epoch Average Loss: 2.059620  
Epoch Average Loss: 2.049451  
Validate Acc: 0.248  
Epoch Average Loss: 2.042102  
Epoch Average Loss: 2.036246  
Epoch Average Loss: 2.032114  
Validate Acc: 0.241  
Epoch Average Loss: 2.028379  
Epoch Average Loss: 2.025600  
Epoch Average Loss: 2.023217  
Validate Acc: 0.249  
Epoch Average Loss: 2.020925  
Epoch Average Loss: 2.019147  
Epoch Average Loss: 2.017291  
Validate Acc: 0.250  
Epoch Average Loss: 2.015686  
Epoch Average Loss: 2.013929  
Epoch Average Loss: 2.012463  
Validate Acc: 0.267  
Epoch Average Loss: 2.010893  
Epoch Average Loss: 2.009251  
Epoch Average Loss: 2.007897  
Validate Acc: 0.273  
Epoch Average Loss: 2.006437  
Epoch Average Loss: 2.004980  
Epoch Average Loss: 2.003933  
Validate Acc: 0.274  
Epoch Average Loss: 2.002378  
Epoch Average Loss: 2.001054  
Epoch Average Loss: 2.000011  
Validate Acc: 0.284  
Epoch Average Loss: 1.999199  
Epoch Average Loss: 1.997546  
Epoch Average Loss: 1.997077  
Validate Acc: 0.281  
Epoch Average Loss: 1.996216  
Epoch Average Loss: 1.995368  
Epoch Average Loss: 1.994716  
Validate Acc: 0.281  
Epoch Average Loss: 1.994102  
Epoch Average Loss: 1.993045  
Epoch Average Loss: 1.992587  
Validate Acc: 0.283  
Epoch Average Loss: 1.992092  
Epoch Average Loss: 1.991264  
Epoch Average Loss: 1.990941  
Validate Acc: 0.277

Epoch Average Loss: 1.990277  
Epoch Average Loss: 1.990155  
Epoch Average Loss: 1.989498  
Validate Acc: 0.276  
Epoch Average Loss: 1.989053  
Epoch Average Loss: 1.988706  
Epoch Average Loss: 1.988257  
Validate Acc: 0.287  
Epoch Average Loss: 1.988022  
Epoch Average Loss: 1.987579  
Epoch Average Loss: 1.986981  
Validate Acc: 0.278  
Epoch Average Loss: 1.986992  
Epoch Average Loss: 1.986734  
Epoch Average Loss: 1.986265  
Validate Acc: 0.283  
Epoch Average Loss: 1.985763  
Epoch Average Loss: 1.985765  
Epoch Average Loss: 1.985496  
Validate Acc: 0.289  
Epoch Average Loss: 1.985440  
Epoch Average Loss: 1.984914  
Epoch Average Loss: 1.984737  
Validate Acc: 0.284  
Epoch Average Loss: 1.984393  
Epoch Average Loss: 1.984273  
Epoch Average Loss: 1.984093  
Validate Acc: 0.283  
Epoch Average Loss: 1.983952  
Epoch Average Loss: 1.983741  
Epoch Average Loss: 1.983317  
Validate Acc: 0.285  
Epoch Average Loss: 1.983105  
Epoch Average Loss: 1.982887  
Epoch Average Loss: 1.982779  
Validate Acc: 0.292  
Epoch Average Loss: 1.982501  
Epoch Average Loss: 1.982100  
Epoch Average Loss: 1.982103  
Validate Acc: 0.290  
Epoch Average Loss: 1.981793  
Epoch Average Loss: 1.981858  
Epoch Average Loss: 1.981502  
Validate Acc: 0.281  
Epoch Average Loss: 1.981714  
Epoch Average Loss: 1.981302  
Epoch Average Loss: 1.981052  
Validate Acc: 0.292  
Epoch Average Loss: 1.981034  
Epoch Average Loss: 1.980459  
Epoch Average Loss: 1.980729  
Validate Acc: 0.290  
Epoch Average Loss: 1.980478  
Epoch Average Loss: 1.980213  
Epoch Average Loss: 1.980048  
Validate Acc: 0.295  
Epoch Average Loss: 1.979646  
Epoch Average Loss: 1.979818  
Epoch Average Loss: 1.979648  
Validate Acc: 0.289  
Epoch Average Loss: 1.979490  
Epoch Average Loss: 1.979336  
Epoch Average Loss: 1.978885  
Validate Acc: 0.285  
Epoch Average Loss: 1.978991

Epoch Average Loss: 1.978719  
Epoch Average Loss: 1.978809  
Validate Acc: 0.289  
Epoch Average Loss: 1.978311  
Epoch Average Loss: 1.978464  
Epoch Average Loss: 1.978114  
Validate Acc: 0.291  
Epoch Average Loss: 1.978040  
Epoch Average Loss: 1.978101  
Epoch Average Loss: 1.977876  
Validate Acc: 0.292  
Epoch Average Loss: 1.977934  
Epoch Average Loss: 1.977653  
Epoch Average Loss: 1.977775  
Validate Acc: 0.292  
Epoch Average Loss: 1.977493  
Epoch Average Loss: 1.977373  
Epoch Average Loss: 1.977231  
Validate Acc: 0.293  
Epoch Average Loss: 1.977331  
Epoch Average Loss: 1.977168  
Epoch Average Loss: 1.976825  
Validate Acc: 0.290  
Epoch Average Loss: 1.976985  
Epoch Average Loss: 1.976800  
Epoch Average Loss: 1.976757  
Validate Acc: 0.294  
Epoch Average Loss: 1.976588  
Epoch Average Loss: 1.976476  
Epoch Average Loss: 1.976556  
Validate Acc: 0.299  
Epoch Average Loss: 1.976479  
Epoch Average Loss: 1.976542  
Epoch Average Loss: 1.976241  
Validate Acc: 0.294  
Epoch Average Loss: 1.976137  
Epoch Average Loss: 1.976267  
Epoch Average Loss: 1.976021  
Validate Acc: 0.300  
Epoch Average Loss: 1.975883  
Epoch Average Loss: 1.975941  
Epoch Average Loss: 1.975911  
Validate Acc: 0.292  
Epoch Average Loss: 1.975781  
Epoch Average Loss: 1.975710  
Epoch Average Loss: 1.975707  
Validate Acc: 0.297  
Epoch Average Loss: 1.975618  
Epoch Average Loss: 1.975506  
Epoch Average Loss: 1.975346  
Validate Acc: 0.293  
Epoch Average Loss: 1.975455  
Epoch Average Loss: 1.975448  
Epoch Average Loss: 1.975436  
Validate Acc: 0.291  
Epoch Average Loss: 1.975223  
Epoch Average Loss: 1.975109  
Epoch Average Loss: 1.975091  
Validate Acc: 0.298  
Epoch Average Loss: 1.975144  
Epoch Average Loss: 1.974871  
Epoch Average Loss: 1.974973  
Validate Acc: 0.297  
Epoch Average Loss: 1.974881  
Epoch Average Loss: 1.974986

Epoch Average Loss: 1.974834  
Validate Acc: 0.291  
Epoch Average Loss: 1.975150  
Epoch Average Loss: 1.974923  
Epoch Average Loss: 1.974797  
Validate Acc: 0.292  
Epoch Average Loss: 1.974650  
Epoch Average Loss: 1.974616  
Epoch Average Loss: 1.974848  
Validate Acc: 0.295  
Epoch Average Loss: 1.974556  
Epoch Average Loss: 1.974466  
Epoch Average Loss: 1.974275  
Validate Acc: 0.290  
Epoch Average Loss: 1.974333  
Epoch Average Loss: 1.974299  
Epoch Average Loss: 1.974518  
Validate Acc: 0.297  
Epoch Average Loss: 1.974331  
Epoch Average Loss: 1.974413  
Epoch Average Loss: 1.974176  
Validate Acc: 0.295  
Epoch Average Loss: 1.974108  
Epoch Average Loss: 1.974335  
Epoch Average Loss: 1.973932  
Validate Acc: 0.289  
Epoch Average Loss: 1.974008  
Epoch Average Loss: 1.974124  
Epoch Average Loss: 1.973978  
Validate Acc: 0.294  
Epoch Average Loss: 1.974023  
Epoch Average Loss: 1.973891  
Epoch Average Loss: 1.973815  
Validate Acc: 0.292  
Epoch Average Loss: 1.974074  
Epoch Average Loss: 1.973775  
Epoch Average Loss: 1.973617  
Validate Acc: 0.294  
Epoch Average Loss: 1.973553  
Epoch Average Loss: 1.973463  
Epoch Average Loss: 1.973590  
Validate Acc: 0.292  
Epoch Average Loss: 1.973279  
Epoch Average Loss: 1.973719  
Epoch Average Loss: 1.973645  
Validate Acc: 0.294  
Epoch Average Loss: 1.973282  
Epoch Average Loss: 1.973418  
Epoch Average Loss: 1.973407  
Validate Acc: 0.292  
Epoch Average Loss: 1.973107  
Epoch Average Loss: 1.973291  
Epoch Average Loss: 1.973255  
Validate Acc: 0.295  
Epoch Average Loss: 1.973200  
Epoch Average Loss: 1.973248  
Epoch Average Loss: 1.973392  
Validate Acc: 0.287  
Epoch Average Loss: 1.973056  
Epoch Average Loss: 1.973167  
Epoch Average Loss: 1.972623  
Validate Acc: 0.289  
Epoch Average Loss: 1.973039  
Epoch Average Loss: 1.973138  
Epoch Average Loss: 1.972953

Validate Acc: 0.293  
Epoch Average Loss: 1.972811  
Epoch Average Loss: 1.972845  
Epoch Average Loss: 1.972635  
Validate Acc: 0.298  
Epoch Average Loss: 1.972745  
Epoch Average Loss: 1.972780  
Epoch Average Loss: 1.972543  
Validate Acc: 0.291  
Epoch Average Loss: 1.972371  
Epoch Average Loss: 1.972634  
Epoch Average Loss: 1.972544  
Validate Acc: 0.293  
Epoch Average Loss: 1.972533  
Epoch Average Loss: 1.972418  
Epoch Average Loss: 1.972067  
Validate Acc: 0.288  
Epoch Average Loss: 1.972675  
Epoch Average Loss: 1.972419  
Epoch Average Loss: 1.972212  
Validate Acc: 0.295  
Epoch Average Loss: 1.972405  
Epoch Average Loss: 1.972135  
Epoch Average Loss: 1.972309  
Validate Acc: 0.292  
Epoch Average Loss: 1.972150  
Epoch Average Loss: 1.972098  
Epoch Average Loss: 1.972139  
Validate Acc: 0.286  
Epoch Average Loss: 1.972155  
Epoch Average Loss: 1.972090  
Epoch Average Loss: 1.972018  
Validate Acc: 0.292  
Epoch Average Loss: 1.971951  
Epoch Average Loss: 1.972149  
Epoch Average Loss: 1.971988  
Validate Acc: 0.296  
Epoch Average Loss: 1.971929  
Epoch Average Loss: 1.971999  
Epoch Average Loss: 1.971938  
Validate Acc: 0.293  
Epoch Average Loss: 1.971797  
Epoch Average Loss: 1.971836  
Epoch Average Loss: 1.971799  
Validate Acc: 0.292  
Epoch Average Loss: 1.971718  
Epoch Average Loss: 1.971439  
Epoch Average Loss: 1.971948  
Validate Acc: 0.289  
Epoch Average Loss: 1.971414  
Epoch Average Loss: 1.971632  
Epoch Average Loss: 1.971705  
Validate Acc: 0.301  
Epoch Average Loss: 1.971485  
Epoch Average Loss: 1.971376  
Epoch Average Loss: 1.971467  
Validate Acc: 0.289  
Epoch Average Loss: 1.971522  
Epoch Average Loss: 1.971554  
Epoch Average Loss: 1.971376  
Validate Acc: 0.295  
Epoch Average Loss: 1.971323  
Epoch Average Loss: 1.971473  
Epoch Average Loss: 1.971461  
Validate Acc: 0.293

```
Epoch Average Loss: 1.971204
Epoch Average Loss: 1.971091
Epoch Average Loss: 1.971100
Validate Acc: 0.297
Epoch Average Loss: 1.971254
Epoch Average Loss: 1.971071
Epoch Average Loss: 1.971188
Validate Acc: 0.290
Epoch Average Loss: 1.970850
Epoch Average Loss: 1.971307
Epoch Average Loss: 1.971074
Validate Acc: 0.296
Epoch Average Loss: 1.970825
Epoch Average Loss: 1.970910
```

## Print the training and validation accuracies for the default hyper-parameters provided

```
In [9]: from utils.evaluation import get_classification_accuracy
out_train = net.predict(x_train)
acc = get_classification_accuracy(out_train, y_train)
print("Training acc: ", acc)
out_val = net.predict(x_val)
acc = get_classification_accuracy(out_val, y_val)
print("Validation acc: ", acc)
```

Training acc: 0.31024444444444443  
Validation acc: 0.295

## Debug the training

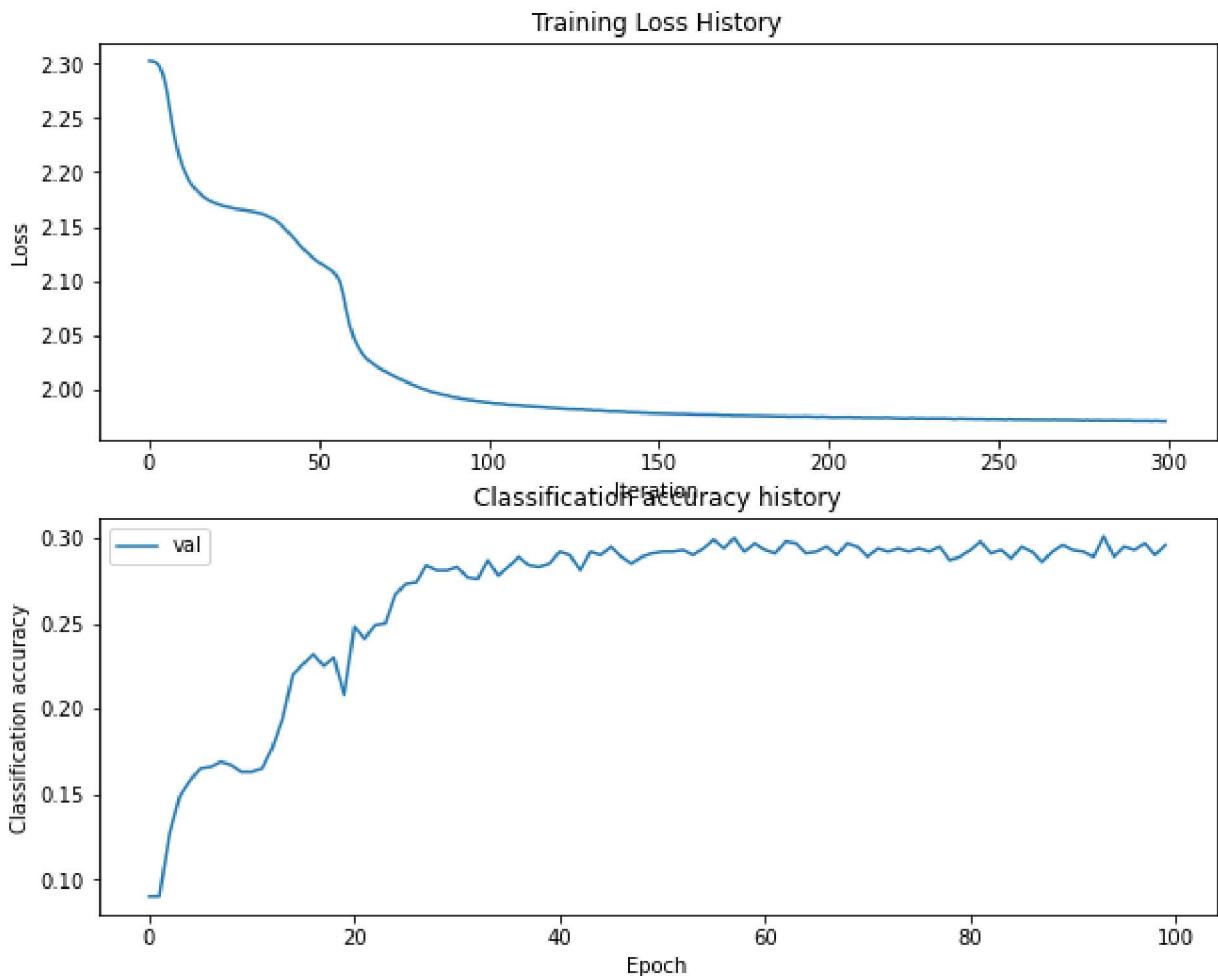
With the default parameters we provided above, you should get a validation accuracy of about 0.29 on the validation set. This isn't very good.

One strategy for getting insight into what's wrong is to plot the training loss function and the validation accuracies during optimization.

Another strategy is to visualize the weights that were learned in the first layer of the network. In most neural networks trained on visual data, the first layer weights typically show some visible structure when visualized.

```
In [10]: # Plot the training Loss function and validation accuracies
plt.subplot(2, 1, 1)
plt.plot(train_error)
plt.title('Training Loss History')
plt.xlabel('Iteration')
plt.ylabel('Loss')

plt.subplot(2, 1, 2)
# plt.plot(stats['train_acc_history'], label='train')
plt.plot(validation_accuracy, label='val')
plt.title('Classification accuracy history')
plt.xlabel('Epoch')
plt.ylabel('Classification accuracy')
plt.legend()
plt.show()
```



```
In [11]: pip install past
```

Note: you may need to restart the kernel to use updated packages.

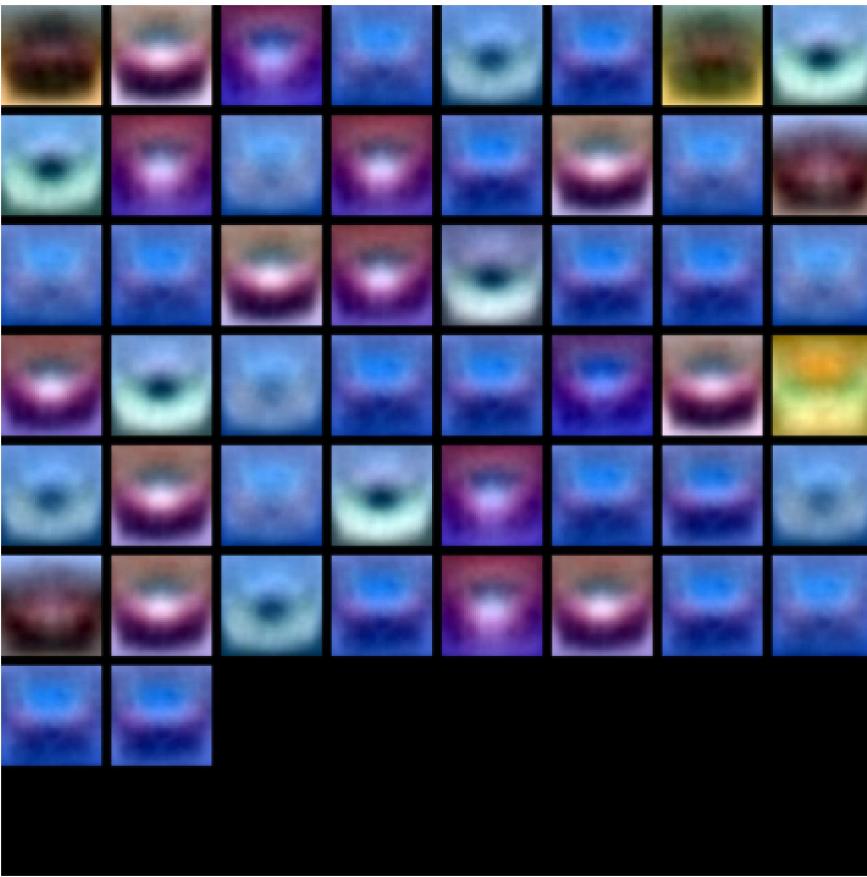
ERROR: Could not find a version that satisfies the requirement past (from versions: none)

ERROR: No matching distribution found for past

```
In [12]: from utils.vis_utils import visualize_grid
# Credits: http://cs231n.stanford.edu/
```

```
# Visualize the weights of the network
def show_net_weights(net):
    W1 = net._modules[0].parameters[0]
    W1 = W1.reshape(3, 32, 32, -1).transpose(3, 1, 2, 0)
    plt.imshow(visualize_grid(W1, padding=3).astype('uint8'))
    plt.gca().axis('off')
    plt.show()

show_net_weights(net)
```



## Tune your hyperparameters (50%)

**What's wrong?** Looking at the visualizations above, we see that the loss is decreasing more or less linearly, which seems to suggest that the learning rate may be too low. Moreover, there is no gap between the training and validation accuracy, suggesting that the model we used has low capacity, and that we should increase its size. On the other hand, with a very large model we would expect to see more overfitting, which would manifest itself as a very large gap between the training and validation accuracy.

**Tuning.** Tuning the hyperparameters and developing intuition for how they affect the final performance is a large part of using Neural Networks, so we want you to get a lot of practice. Below, you should experiment with different values of the various hyperparameters, including hidden layer size, learning rate, numer of training epochs, and regularization strength.

**Approximate results.** You should be aim to achieve a classification accuracy of greater than 48% on the validation set. Our best network gets over 52% on the validation set.

**Experiment:** Your goal in this exercise is to get as good of a result on CIFAR-10 as you can (52% could serve as a reference), with a fully-connected Neural Network.

**Explain your hyperparameter tuning process below.**

*Your Answer :*

I upped the learning rate to 0.01 and got an accuracy below 52% so I upped to 0.025 and upped the weight decay to 0.001 but the accuracy was still low so I put it to 0.0001 and with the 0.025 learning rate and got the accuracy I wanted

```
In [13]: best_net_hyperparams = [0.025, 0.0001, 300, 100] # store the best model into this

#####
# TODO: Tune hyperparameters using the validation set. Store your best trained #
# model hyperparams in best_net.
#
# To help debug your network, it may help to use visualizations similar to the #
# ones we used above; these visualizations will have significant qualitative #
# differences from the ones we saw above for the poorly tuned network.
#
# You are now free to test different combinations of hyperparameters to build #
# various models and test them according to the above plots and visualization #

#
# TODO: Show the above plots and visualizations for the default params (already #
# done) and the best hyper-params you obtain. You only need to show this for 2 #
# sets of hyper-params.
#
# You just need to store values for the hyperparameters in best_net_hyperparams #
# as a List in the order
# best_net_hyperparams = [lr, weight_decay, epoch, hidden_size]
#####

Lr = best_net_hyperparams[0]
Weightdecay = best_net_hyperparams[1]
epochs = best_net_hyperparams[2]
batch_size = best_net_hyperparams[3]

nu_net = init_model()

optim = SGD(nu_net, lr=Lr, weight_decay=Weightdecay)
loss_func = CrossEntropyLoss()

trainer = Trainer(dataset, optim, nu_net, loss_func, epochs, batch_size, validate_inter
train_error, validation_accuracy = trainer.train()

Epoch Average Loss: 2.241096
Validate Acc: 0.169
Epoch Average Loss: 2.125412
Epoch Average Loss: 2.036844
Epoch Average Loss: 1.955349
Validate Acc: 0.287
Epoch Average Loss: 1.895586
Epoch Average Loss: 1.848127
Epoch Average Loss: 1.806558
Validate Acc: 0.329
Epoch Average Loss: 1.766173
Epoch Average Loss: 1.731965
Epoch Average Loss: 1.705697
Validate Acc: 0.396
Epoch Average Loss: 1.680954
Epoch Average Loss: 1.666421
Epoch Average Loss: 1.647634
Validate Acc: 0.416
Epoch Average Loss: 1.631943
Epoch Average Loss: 1.618325
Epoch Average Loss: 1.604315
Validate Acc: 0.408
```

Epoch Average Loss: 1.595931  
Epoch Average Loss: 1.584065  
Epoch Average Loss: 1.576126  
Validate Acc: 0.443  
Epoch Average Loss: 1.562295  
Epoch Average Loss: 1.551990  
Epoch Average Loss: 1.544662  
Validate Acc: 0.426  
Epoch Average Loss: 1.539534  
Epoch Average Loss: 1.530122  
Epoch Average Loss: 1.526384  
Validate Acc: 0.454  
Epoch Average Loss: 1.511151  
Epoch Average Loss: 1.513174  
Epoch Average Loss: 1.504139  
Validate Acc: 0.436  
Epoch Average Loss: 1.497620  
Epoch Average Loss: 1.494806  
Epoch Average Loss: 1.493054  
Validate Acc: 0.472  
Epoch Average Loss: 1.484712  
Epoch Average Loss: 1.482217  
Epoch Average Loss: 1.476422  
Validate Acc: 0.454  
Epoch Average Loss: 1.473267  
Epoch Average Loss: 1.470186  
Epoch Average Loss: 1.468078  
Validate Acc: 0.453  
Epoch Average Loss: 1.467035  
Epoch Average Loss: 1.463855  
Epoch Average Loss: 1.459286  
Validate Acc: 0.462  
Epoch Average Loss: 1.457692  
Epoch Average Loss: 1.454057  
Epoch Average Loss: 1.452056  
Validate Acc: 0.484  
Epoch Average Loss: 1.450131  
Epoch Average Loss: 1.447299  
Epoch Average Loss: 1.444539  
Validate Acc: 0.463  
Epoch Average Loss: 1.445107  
Epoch Average Loss: 1.438764  
Epoch Average Loss: 1.438547  
Validate Acc: 0.477  
Epoch Average Loss: 1.437946  
Epoch Average Loss: 1.437197  
Epoch Average Loss: 1.433660  
Validate Acc: 0.482  
Epoch Average Loss: 1.429819  
Epoch Average Loss: 1.433175  
Epoch Average Loss: 1.428987  
Validate Acc: 0.452  
Epoch Average Loss: 1.429116  
Epoch Average Loss: 1.424877  
Epoch Average Loss: 1.423645  
Validate Acc: 0.487  
Epoch Average Loss: 1.424995  
Epoch Average Loss: 1.419581  
Epoch Average Loss: 1.420101  
Validate Acc: 0.470  
Epoch Average Loss: 1.419877  
Epoch Average Loss: 1.415901  
Epoch Average Loss: 1.415233  
Validate Acc: 0.487  
Epoch Average Loss: 1.409624

Epoch Average Loss: 1.415624  
Epoch Average Loss: 1.409490  
Validate Acc: 0.486  
Epoch Average Loss: 1.407335  
Epoch Average Loss: 1.411262  
Epoch Average Loss: 1.407392  
Validate Acc: 0.475  
Epoch Average Loss: 1.407123  
Epoch Average Loss: 1.402580  
Epoch Average Loss: 1.405636  
Validate Acc: 0.474  
Epoch Average Loss: 1.399450  
Epoch Average Loss: 1.398835  
Epoch Average Loss: 1.399392  
Validate Acc: 0.467  
Epoch Average Loss: 1.398352  
Epoch Average Loss: 1.402007  
Epoch Average Loss: 1.397436  
Validate Acc: 0.475  
Epoch Average Loss: 1.397347  
Epoch Average Loss: 1.392460  
Epoch Average Loss: 1.399742  
Validate Acc: 0.475  
Epoch Average Loss: 1.389381  
Epoch Average Loss: 1.391759  
Epoch Average Loss: 1.395754  
Validate Acc: 0.499  
Epoch Average Loss: 1.390833  
Epoch Average Loss: 1.394161  
Epoch Average Loss: 1.390992  
Validate Acc: 0.514  
Epoch Average Loss: 1.381151  
Epoch Average Loss: 1.390230  
Epoch Average Loss: 1.384808  
Validate Acc: 0.470  
Epoch Average Loss: 1.390014  
Epoch Average Loss: 1.387721  
Epoch Average Loss: 1.383815  
Validate Acc: 0.473  
Epoch Average Loss: 1.379634  
Epoch Average Loss: 1.382199  
Epoch Average Loss: 1.379247  
Validate Acc: 0.489  
Epoch Average Loss: 1.382114  
Epoch Average Loss: 1.378393  
Epoch Average Loss: 1.378804  
Validate Acc: 0.500  
Epoch Average Loss: 1.385585  
Epoch Average Loss: 1.376533  
Epoch Average Loss: 1.377670  
Validate Acc: 0.507  
Epoch Average Loss: 1.376171  
Epoch Average Loss: 1.377519  
Epoch Average Loss: 1.375801  
Validate Acc: 0.508  
Epoch Average Loss: 1.371977  
Epoch Average Loss: 1.374274  
Epoch Average Loss: 1.377641  
Validate Acc: 0.498  
Epoch Average Loss: 1.376767  
Epoch Average Loss: 1.370655  
Epoch Average Loss: 1.373555  
Validate Acc: 0.503  
Epoch Average Loss: 1.368691  
Epoch Average Loss: 1.370300

Epoch Average Loss: 1.374292  
Validate Acc: 0.499  
Epoch Average Loss: 1.372288  
Epoch Average Loss: 1.371779  
Epoch Average Loss: 1.374580  
Validate Acc: 0.456  
Epoch Average Loss: 1.370122  
Epoch Average Loss: 1.369882  
Epoch Average Loss: 1.374022  
Validate Acc: 0.475  
Epoch Average Loss: 1.370188  
Epoch Average Loss: 1.368587  
Epoch Average Loss: 1.372111  
Validate Acc: 0.512  
Epoch Average Loss: 1.369422  
Epoch Average Loss: 1.366030  
Epoch Average Loss: 1.368338  
Validate Acc: 0.505  
Epoch Average Loss: 1.365068  
Epoch Average Loss: 1.367221  
Epoch Average Loss: 1.368972  
Validate Acc: 0.462  
Epoch Average Loss: 1.363119  
Epoch Average Loss: 1.363624  
Epoch Average Loss: 1.367097  
Validate Acc: 0.521  
Epoch Average Loss: 1.369830  
Epoch Average Loss: 1.366397  
Epoch Average Loss: 1.363651  
Validate Acc: 0.497  
Epoch Average Loss: 1.364659  
Epoch Average Loss: 1.365223  
Epoch Average Loss: 1.362226  
Validate Acc: 0.503  
Epoch Average Loss: 1.362969  
Epoch Average Loss: 1.369488  
Epoch Average Loss: 1.365423  
Validate Acc: 0.491  
Epoch Average Loss: 1.360710  
Epoch Average Loss: 1.358946  
Epoch Average Loss: 1.364710  
Validate Acc: 0.483  
Epoch Average Loss: 1.361687  
Epoch Average Loss: 1.362754  
Epoch Average Loss: 1.363598  
Validate Acc: 0.499  
Epoch Average Loss: 1.362949  
Epoch Average Loss: 1.358705  
Epoch Average Loss: 1.360099  
Validate Acc: 0.516  
Epoch Average Loss: 1.360193  
Epoch Average Loss: 1.361180  
Epoch Average Loss: 1.363828  
Validate Acc: 0.491  
Epoch Average Loss: 1.360838  
Epoch Average Loss: 1.363654  
Epoch Average Loss: 1.356647  
Validate Acc: 0.506  
Epoch Average Loss: 1.361900  
Epoch Average Loss: 1.359774  
Epoch Average Loss: 1.362647  
Validate Acc: 0.496  
Epoch Average Loss: 1.358857  
Epoch Average Loss: 1.356474  
Epoch Average Loss: 1.357523

Validate Acc: 0.500  
Epoch Average Loss: 1.361581  
Epoch Average Loss: 1.358316  
Epoch Average Loss: 1.360654  
Validate Acc: 0.476  
Epoch Average Loss: 1.356717  
Epoch Average Loss: 1.358212  
Epoch Average Loss: 1.355439  
Validate Acc: 0.502  
Epoch Average Loss: 1.359548  
Epoch Average Loss: 1.359769  
Epoch Average Loss: 1.357440  
Validate Acc: 0.507  
Epoch Average Loss: 1.358119  
Epoch Average Loss: 1.355149  
Epoch Average Loss: 1.358511  
Validate Acc: 0.495  
Epoch Average Loss: 1.357007  
Epoch Average Loss: 1.357059  
Epoch Average Loss: 1.355231  
Validate Acc: 0.512  
Epoch Average Loss: 1.357311  
Epoch Average Loss: 1.358459  
Epoch Average Loss: 1.357125  
Validate Acc: 0.494  
Epoch Average Loss: 1.354064  
Epoch Average Loss: 1.357340  
Epoch Average Loss: 1.353602  
Validate Acc: 0.495  
Epoch Average Loss: 1.356047  
Epoch Average Loss: 1.352962  
Epoch Average Loss: 1.353475  
Validate Acc: 0.460  
Epoch Average Loss: 1.354139  
Epoch Average Loss: 1.352252  
Epoch Average Loss: 1.353474  
Validate Acc: 0.498  
Epoch Average Loss: 1.356667  
Epoch Average Loss: 1.355219  
Epoch Average Loss: 1.354171  
Validate Acc: 0.475  
Epoch Average Loss: 1.354713  
Epoch Average Loss: 1.353511  
Epoch Average Loss: 1.352644  
Validate Acc: 0.491  
Epoch Average Loss: 1.355139  
Epoch Average Loss: 1.354011  
Epoch Average Loss: 1.355600  
Validate Acc: 0.486  
Epoch Average Loss: 1.353374  
Epoch Average Loss: 1.352596  
Epoch Average Loss: 1.351523  
Validate Acc: 0.503  
Epoch Average Loss: 1.354047  
Epoch Average Loss: 1.351482  
Epoch Average Loss: 1.352086  
Validate Acc: 0.504  
Epoch Average Loss: 1.351442  
Epoch Average Loss: 1.349260  
Epoch Average Loss: 1.351557  
Validate Acc: 0.488  
Epoch Average Loss: 1.351664  
Epoch Average Loss: 1.347298  
Epoch Average Loss: 1.352837  
Validate Acc: 0.476

Epoch Average Loss: 1.352573  
Epoch Average Loss: 1.348850  
Epoch Average Loss: 1.350201  
Validate Acc: 0.507  
Epoch Average Loss: 1.349109  
Epoch Average Loss: 1.355248  
Epoch Average Loss: 1.353443  
Validate Acc: 0.477  
Epoch Average Loss: 1.351224  
Epoch Average Loss: 1.349944  
Epoch Average Loss: 1.345762  
Validate Acc: 0.510  
Epoch Average Loss: 1.352956  
Epoch Average Loss: 1.356812  
Epoch Average Loss: 1.349443  
Validate Acc: 0.531  
Epoch Average Loss: 1.349989  
Epoch Average Loss: 1.353487  
Epoch Average Loss: 1.348763  
Validate Acc: 0.496  
Epoch Average Loss: 1.349357  
Epoch Average Loss: 1.352284  
Epoch Average Loss: 1.349153  
Validate Acc: 0.488  
Epoch Average Loss: 1.346898  
Epoch Average Loss: 1.351128  
Epoch Average Loss: 1.349654  
Validate Acc: 0.490  
Epoch Average Loss: 1.351600  
Epoch Average Loss: 1.348804  
Epoch Average Loss: 1.350393  
Validate Acc: 0.471  
Epoch Average Loss: 1.347717  
Epoch Average Loss: 1.350751  
Epoch Average Loss: 1.349475  
Validate Acc: 0.505  
Epoch Average Loss: 1.353384  
Epoch Average Loss: 1.351796  
Epoch Average Loss: 1.350365  
Validate Acc: 0.490  
Epoch Average Loss: 1.349336  
Epoch Average Loss: 1.351403  
Epoch Average Loss: 1.348818  
Validate Acc: 0.516  
Epoch Average Loss: 1.348049  
Epoch Average Loss: 1.352009  
Epoch Average Loss: 1.347701  
Validate Acc: 0.504  
Epoch Average Loss: 1.350013  
Epoch Average Loss: 1.344317  
Epoch Average Loss: 1.348380  
Validate Acc: 0.511  
Epoch Average Loss: 1.349290  
Epoch Average Loss: 1.345754  
Epoch Average Loss: 1.348424  
Validate Acc: 0.504  
Epoch Average Loss: 1.345815  
Epoch Average Loss: 1.352811  
Epoch Average Loss: 1.352443  
Validate Acc: 0.488  
Epoch Average Loss: 1.349481  
Epoch Average Loss: 1.345980  
Epoch Average Loss: 1.343192  
Validate Acc: 0.485  
Epoch Average Loss: 1.346689

```
Epoch Average Loss: 1.349507
Epoch Average Loss: 1.343067
Validate Acc: 0.507
Epoch Average Loss: 1.352967
Epoch Average Loss: 1.350717
Epoch Average Loss: 1.345463
Validate Acc: 0.519
Epoch Average Loss: 1.344334
Epoch Average Loss: 1.344041
Epoch Average Loss: 1.347282
Validate Acc: 0.496
Epoch Average Loss: 1.345986
Epoch Average Loss: 1.343411
Epoch Average Loss: 1.342492
Validate Acc: 0.506
Epoch Average Loss: 1.345104
Epoch Average Loss: 1.347813
Epoch Average Loss: 1.347679
Validate Acc: 0.471
Epoch Average Loss: 1.343445
Epoch Average Loss: 1.345334
Epoch Average Loss: 1.345379
Validate Acc: 0.489
Epoch Average Loss: 1.349347
Epoch Average Loss: 1.341943
Epoch Average Loss: 1.345120
Validate Acc: 0.486
Epoch Average Loss: 1.345382
Epoch Average Loss: 1.347938
Epoch Average Loss: 1.350910
Validate Acc: 0.488
Epoch Average Loss: 1.344121
Epoch Average Loss: 1.346478
Epoch Average Loss: 1.344316
Validate Acc: 0.505
Epoch Average Loss: 1.347673
Epoch Average Loss: 1.344596
Epoch Average Loss: 1.348383
Validate Acc: 0.516
Epoch Average Loss: 1.343126
Epoch Average Loss: 1.342402
Epoch Average Loss: 1.345841
Validate Acc: 0.504
Epoch Average Loss: 1.348359
Epoch Average Loss: 1.343954
Epoch Average Loss: 1.343444
Validate Acc: 0.499
Epoch Average Loss: 1.346912
Epoch Average Loss: 1.342198
Epoch Average Loss: 1.344352
Validate Acc: 0.496
Epoch Average Loss: 1.343037
Epoch Average Loss: 1.343527
```

```
In [14]: # TODO: Plot the training_error and validation_accuracy of the best network (5%)
out_train = nu_net.predict(x_train)
acc = get_classification_accuracy(out_train, y_train)
print("Training acc: ", acc)
out_val = nu_net.predict(x_val)
acc = get_classification_accuracy(out_val, y_val)
print("Validation acc: ", acc)

# Plot the training Loss function and validation accuracies
plt.subplot(2, 1, 1)
```

```

plt.plot(train_error)
plt.title('Training Loss History')
plt.xlabel('Iteration')
plt.ylabel('Loss')

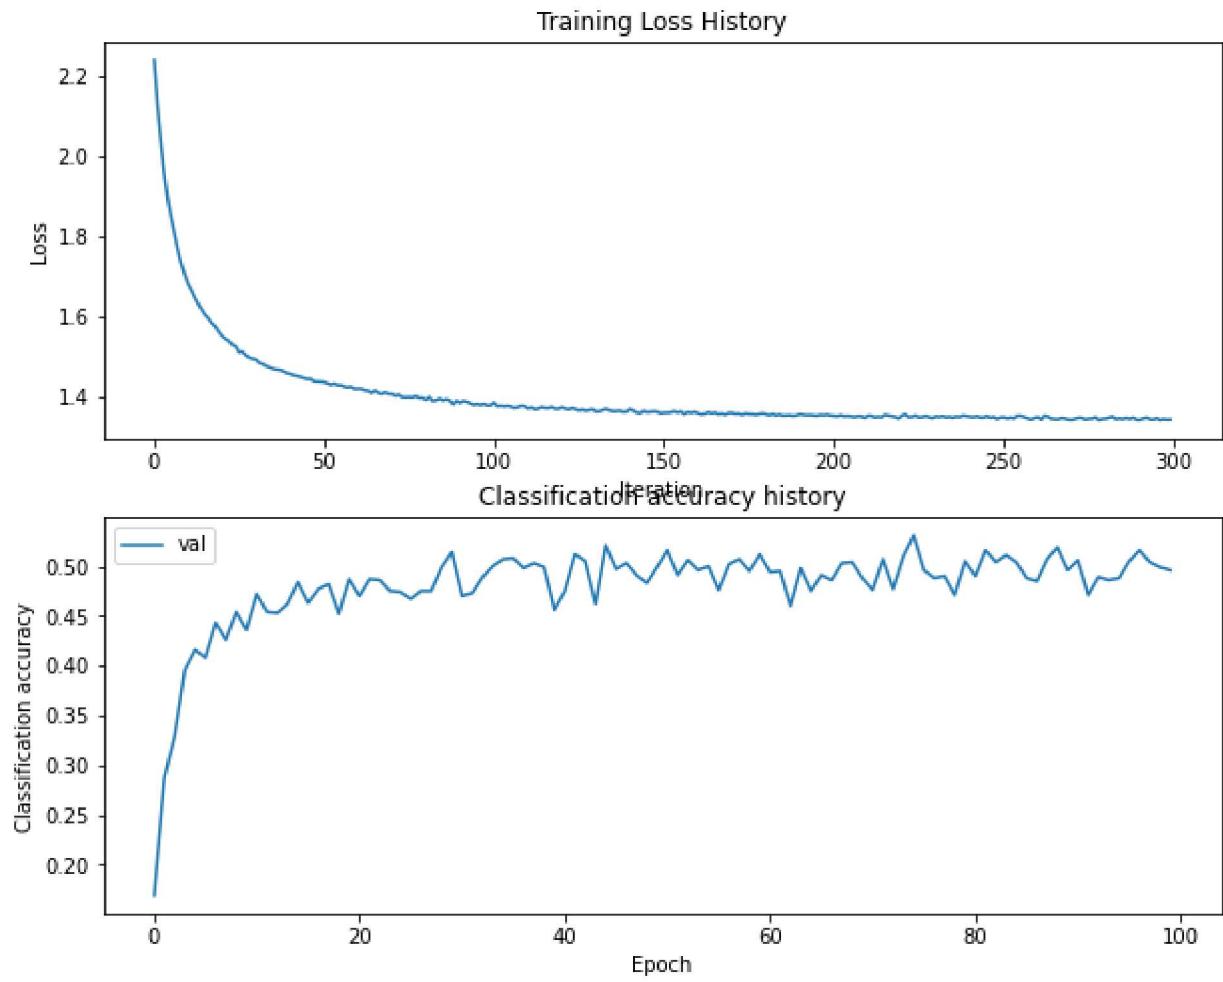
plt.subplot(2, 1, 2)
# plt.plot(stats['train_acc_history'], label='train')
plt.plot(validation_accuracy, label='val')
plt.title('Classification accuracy history')
plt.xlabel('Epoch')
plt.ylabel('Classification accuracy')
plt.legend()
plt.show()

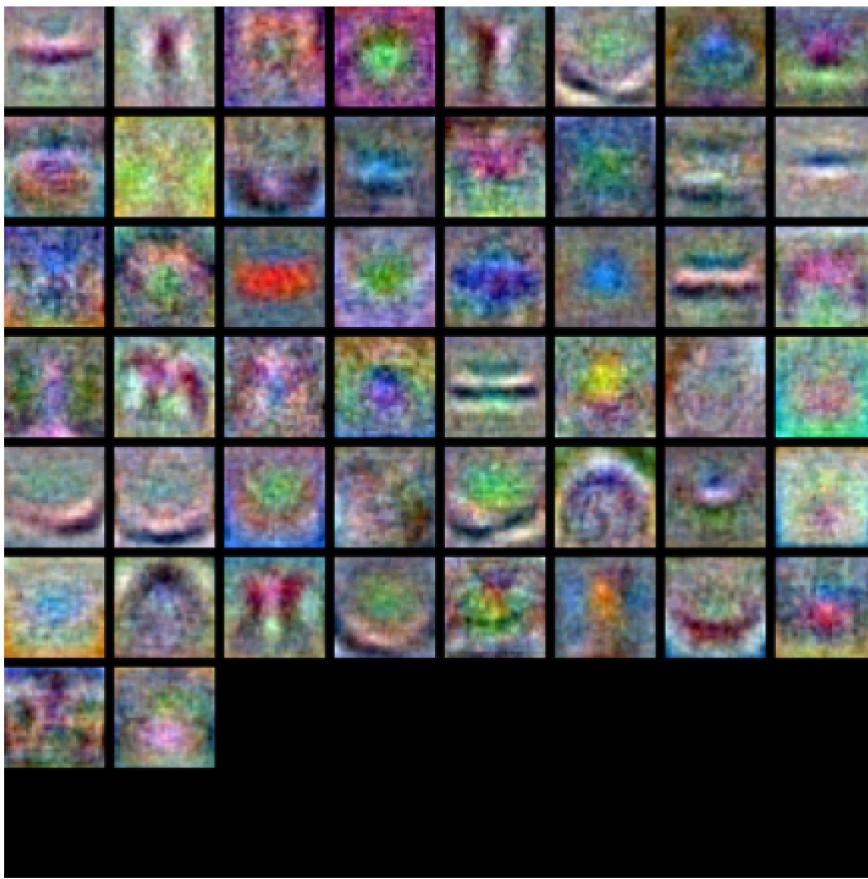
# TODO: visualize the weights of the best network (5%)

show_net_weights(nu_net)

```

Training acc: 0.4857555555555555  
Validation acc: 0.453





## Run on the test set (30%)

When you are done experimenting, you should evaluate your final trained network on the test set; you should get above 48%.

```
In [17]: test_acc = (nu_net.predict(x_test) == y_test).mean()  
print('Test accuracy: ', test_acc)
```

Test accuracy: 0.4483

### Inline Question (10%)

Now that you have trained a Neural Network classifier, you may find that your testing accuracy is much lower than the training accuracy. In what ways can we decrease this gap? Select all that apply.

1. Train on a larger dataset.
2. Add more hidden units.
3. Increase the regularization strength.
4. None of the above.

*Your Answer :* My answer is all 3 of the above would decrease the gap.

*Your Explanation :* Training on a larger dataset, adding more hidden units, and increasing the regularization strength will decrease the gap. Training on a larger dataset will decrease the gap because the model will become more accustomed to a larger variety of data, adding more hidden

units as the model will be able to train to more general data. Increasing regularization strength will reduce overfitting from the model training on training data.

In [ ]: